



2

CS 3104 OPERATING SYSTEMS

CHAPTER 2 OPERATING-SYSTEM STRUCTURES





OUTLINE

2

✓ Outline

- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



WHY APPLICATIONS ARE OPERATING-SYSTEM SPECIFIC?

- Outline
- ✓ Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



- **Fundamental Fact:**
 - Apps compiled on one operating system are not usually executable on other operating systems
- **Problem:**
 - Each operating system provides its own unique system calls
 - Own file formats, etc.



WHY APPLICATIONS ARE OPERATING-SYSTEM SPECIFIC?

- Outline
- ✓ Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



- **Apps can be multi-operating system:**
 1. Apps written in interpreted language (like Python or Ruby) that has an interpreter available on multiple operating systems
 2. Apps written in language that includes a VM containing the running app (like Java)
 3. Application developer can use a standard language (like C) or API, compile separately on each operating system to run on each
- **Application Binary Interface (ABI):**
 - architecture equivalent of API (**API**: application level)
 - defines how different components of binary code can interface for a given operating system on a given architecture, CPU, etc.



WHY APPLICATIONS ARE OPERATING-SYSTEM SPECIFIC?

Important Notes

- Unless an interpreter, RTE, or binary executable file is written for and compiled on a specific operating system on a specific CPU type (such as Intel x86 or ARMv8), the application will fail to run.
- Imagine the amount of work that is required for a program such as the Firefox browser to run on Windows, macOS, various Linux releases, iOS, and Android, sometimes on various CPU architectures.



- Outline
- ✓ Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



OPERATING-SYSTEM DESIGN AND IMPLEMENTATION

- Outline
- Why Applications are Operating-System Specific?
- ✓ Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



- Problems in the design and implementation of OS that are not “solvable”, but some approaches have proven successful
- Internal structure of different Operating Systems can vary widely
- Start the design by defining goals and specifications
- Affected by choice of hardware and type of system
 - traditional desktop/laptop, mobile, distributed, or real time



OPERATING-SYSTEM DESIGN AND IMPLEMENTATION

- Outline
- Why Applications are Operating-System Specific?
- ✓ Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



▪ User Goals:

- operating system should be convenient to use, easy to learn, reliable, safe, and fast

▪ System Goals:

- operating system should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient

▪ Important Notes:

- requirements are vague and may be interpreted in various ways
- no unique solution to the problem of defining the requirements for an operating system
- wide range of systems in existence shows that different requirements can result in a large variety of solutions for different environments



OPERATING-SYSTEM DESIGN AND IMPLEMENTATION

- Outline
- Why Applications are Operating-System Specific?
- ✓ Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



- Important principle to separate:
Policy: *What* will be done?
Mechanism: *How* to do it?
- **Mechanisms** determine how to do something
- **Policies** decide what will be done
- The **separation of policy from mechanism** is a very important principle, it allows maximum flexibility if policy decisions are to be changed later (example: timer)
- **Specifying and designing an OS is highly creative task of software engineering**



OPERATING-SYSTEM DESIGN AND IMPLEMENTATION

OS Implementation

- **Operating Systems are collections of many programs**
 - written by many people over a long period of time
 - difficult to make general statements on their implementation
- **Much variation:**
 - Early Operating Systems were written in Assembly Language
 - Then system programming languages like Algol, PL/1
 - Now in C, C++



- Outline
- Why Applications are Operating-System Specific?
- ✓ Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



OPERATING-SYSTEM DESIGN AND IMPLEMENTATION

OS Implementation

- Actually, usually a **mix of languages**:
 - Lowest levels of the kernel in assembly
 - Main body in C
 - Systems programs in C, C++, scripting languages like PERL, Python, shell scripts
- **More high-level language** are easier to **port** to other hardware
 - but slower
- **Emulation** can allow an OS to run on non-native hardware



- Outline
- Why Applications are Operating-System Specific?
- ✓ Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



OPERATING-SYSTEM DESIGN AND IMPLEMENTATION

Important Notes

- Outline
- Why Applications are Operating-System Specific?
- ✓ Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



- Major performance improvements in operating systems are more likely to be the result of **better data structures and algorithms** than of **excellent assembly-language code**.
- In addition, although operating systems are large, **only a small amount of the code is critical to high performance**; the **interrupt handlers, I/O manager, memory manager, and CPU scheduler** are probably the most critical routines.
- After the system is written and is working correctly, bottlenecks can be identified and can be refactored to operate more efficiently



OPERATING-SYSTEM STRUCTURE

- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



- General-purpose OS is a very large and complex program
- Various ways to structure Operating Systems:
 - **Simple structure:** MS-DOS
 - **More complex:** UNIX
 - **Layered:** An abstraction
 - **Microkernel:** Mach

Note: Not all versions of Mach are microkernels



OPERATING-SYSTEM STRUCTURE

Monolithic Structure

- **Monolithic Structure:**
 - common technique for designing OS
 - place all of the functionality of the kernel into a single, static binary file that runs in a single address space
 - often known as a tightly coupled system because changes to one part of the system can have wide-ranging effects on other parts
 - **example: Original UNIX OS**



- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



OPERATING-SYSTEM STRUCTURE

Monolithic Structure

- **UNIX OS**
 - limited by hardware functionality
 - the original UNIX operating system had limited structuring
 - **Consists of two separable parts:**
 - **Systems programs**
 - **The kernel**
 - ✓ Consists of everything below the system-call interface and above the physical hardware
 - ✓ Provides the file system, CPU scheduling, memory management, and other operating-system functions; a large number of functions for one level

- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging

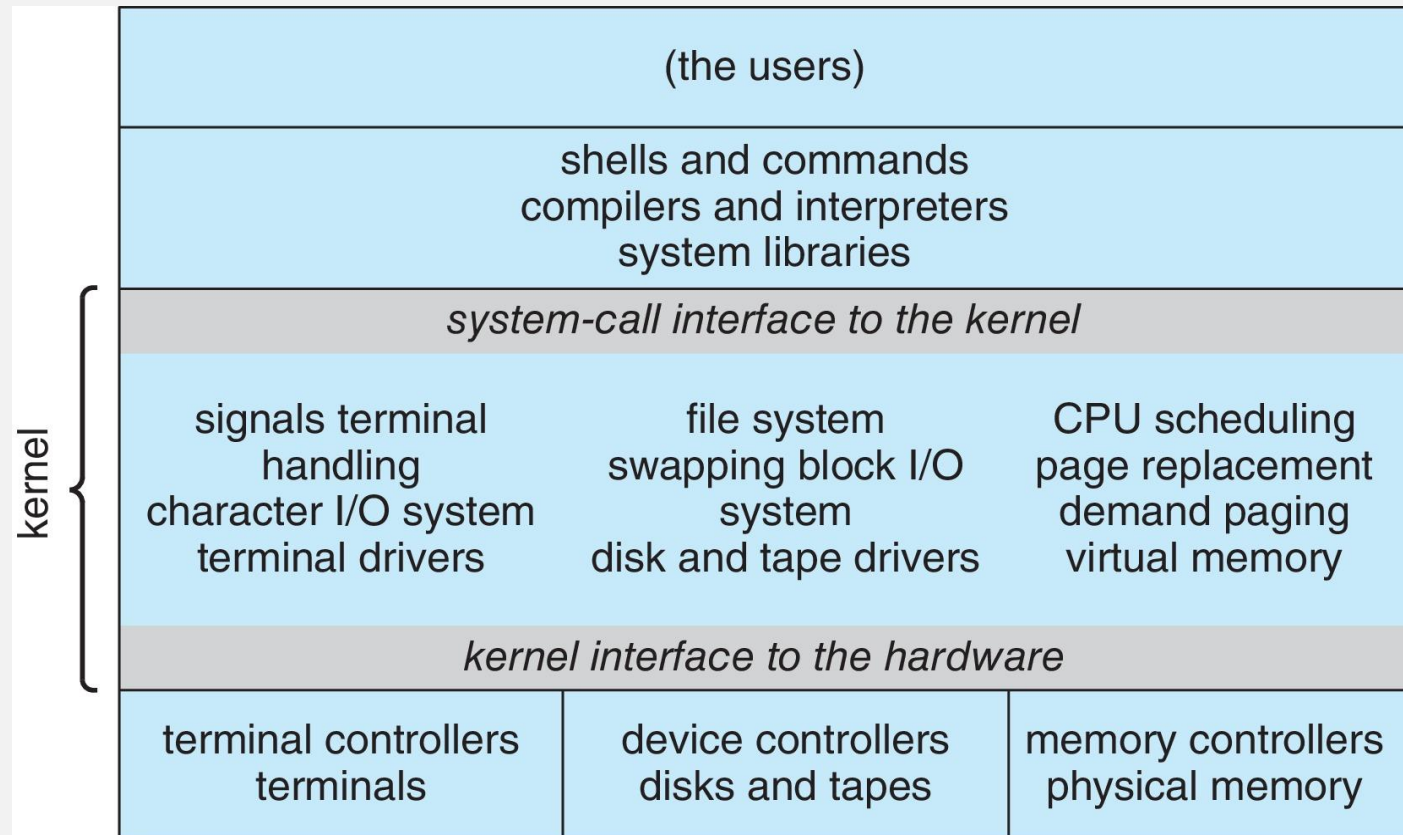




OPERATING-SYSTEM STRUCTURE

Traditional UNIX System Structure

Beyond simple but not fully layered



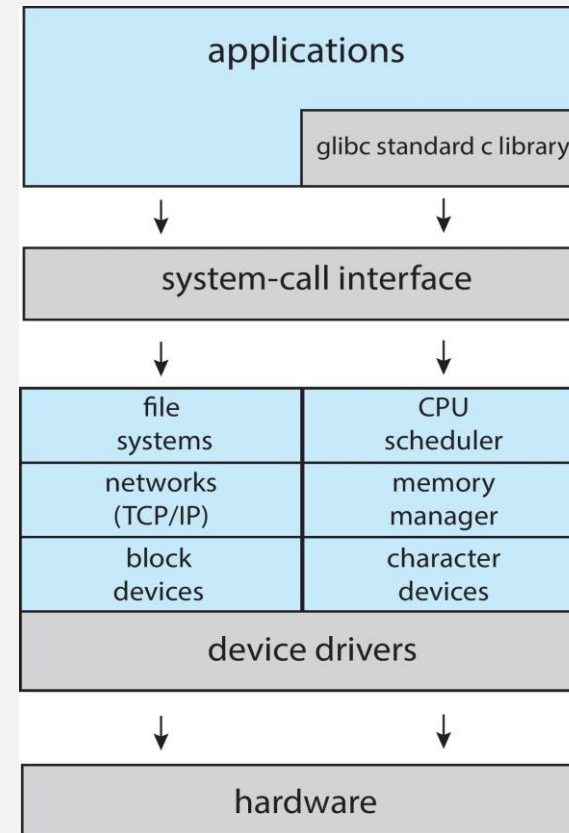
- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



OPERATING-SYSTEM STRUCTURE

Linux System Structure

Monolithic plus modular design



- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging






OPERATING-SYSTEM STRUCTURE

Layered Approach

- The operating system is divided into a number of layers (levels), each is built on top of lower layers.
- The **bottom layer (layer 0)**, is the **hardware**; the **highest (layer N)** is the **user interface**.
- With modularity (**modular approach**), layers are selected such that **each** uses **functions** (operations) and **services** of only lower-level layers



- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



OPERATING-SYSTEM STRUCTURE

Microkernels

- A piece of software or even code that contains the **near-minimum** amount of functions and features required to **implement an** operating system.
- Moves as much **from the kernel** into **user space**
- **Mach**: example of **microkernel**
 - Mac OS X kernel (**Darwin**) is partly based on Mach
- **Communication** takes place between user modules using **message passing**



- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



OPERATING-SYSTEM STRUCTURE

Microkernels

■ Benefits:

- Easier to extend a microkernel
- Easier to port the operating system to new architectures
- More reliable (less code is running in kernel mode)
- More secure

■ Detriments:

- Performance **overhead** of user space to kernel space communication

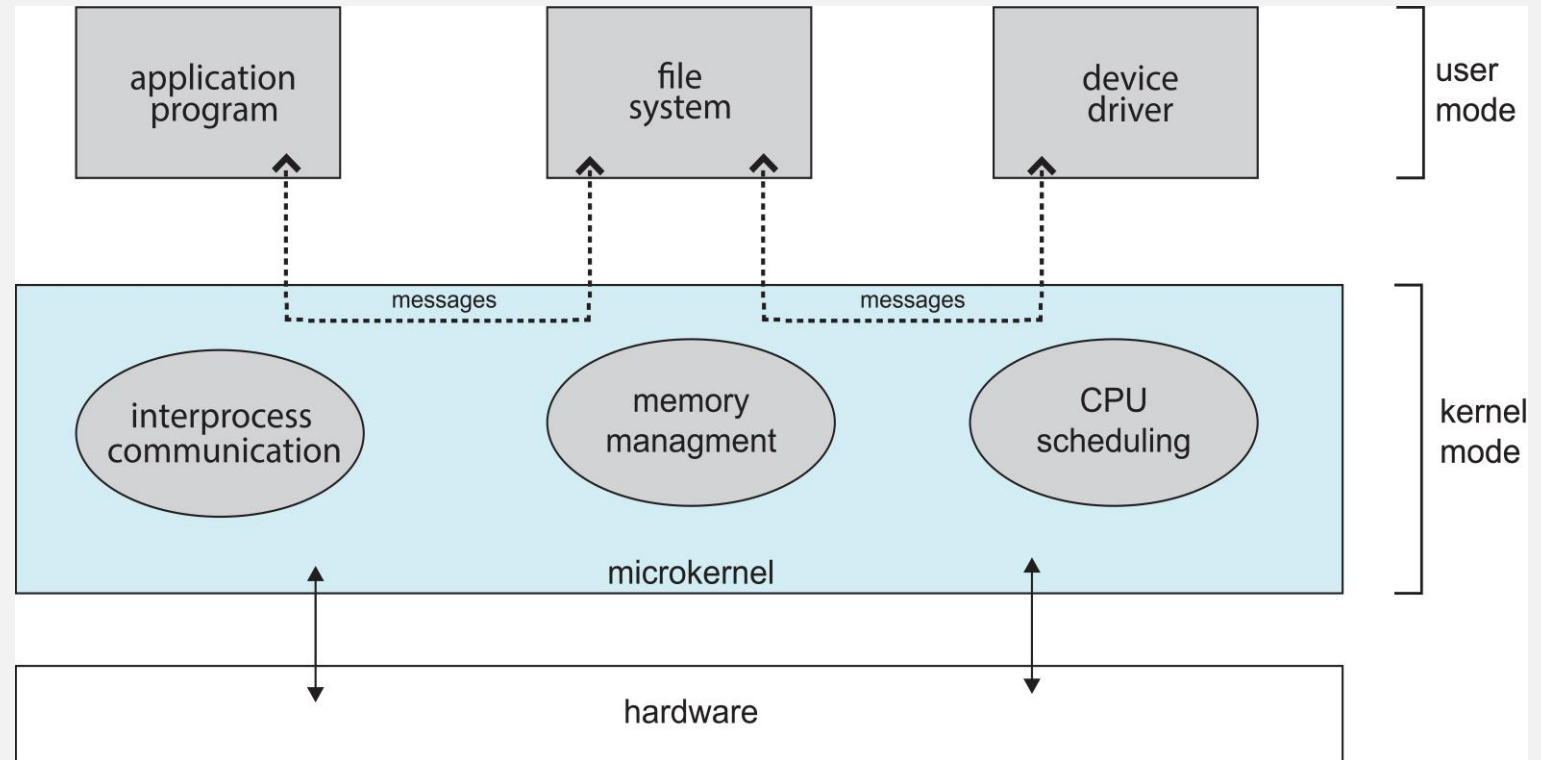


- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



OPERATING-SYSTEM STRUCTURE

Microkernel System Structure



- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



OPERATING-SYSTEM STRUCTURE

Modules

- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging



- Many modern operating systems implement **Loadable Kernel Modules (LKMs)**
 - Uses object-oriented approach
 - Each core component is separate
 - Each talks to the others over known interfaces
 - Each is loadable as needed within the kernel
- Overall, **similar to layered approach** but **more flexible** because **any module can call any other module**
 - Linux, Solaris, etc.



OPERATING-SYSTEM STRUCTURE

Hybrid Systems

- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging

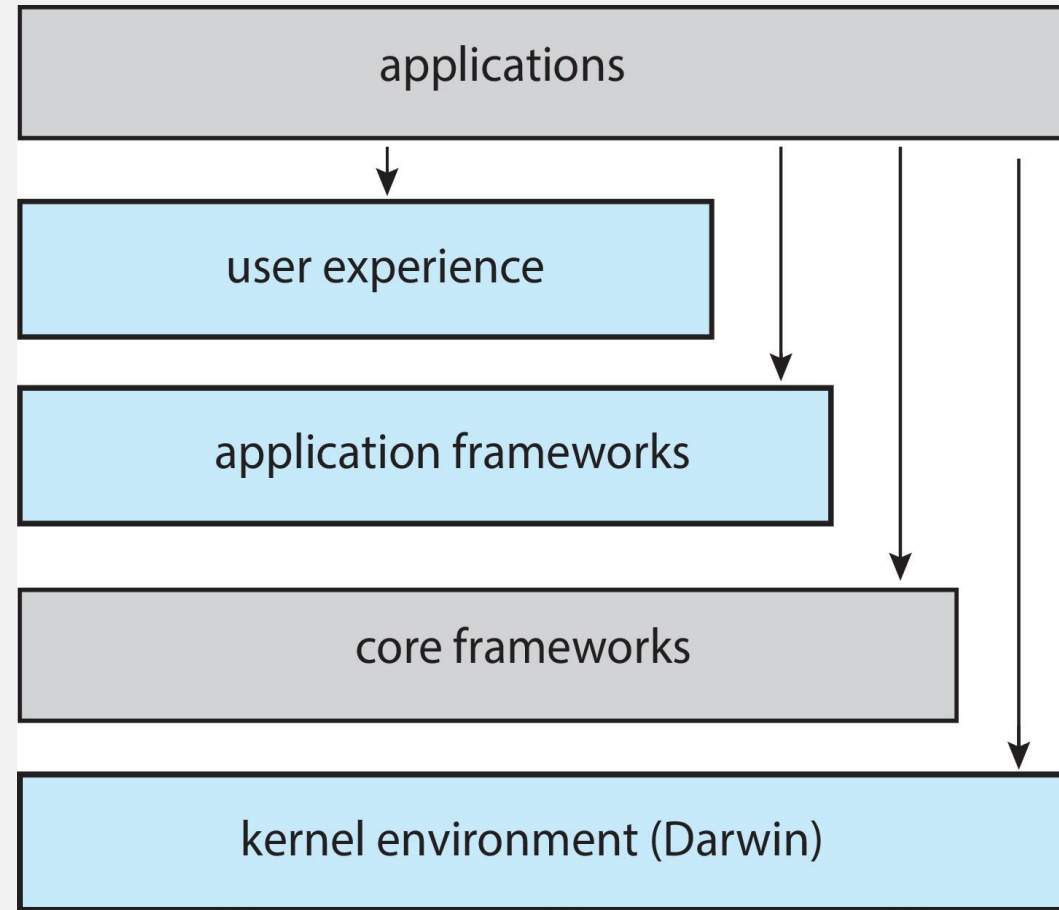


- Most modern operating systems are actually not one pure model
 - Hybrid combines multiple approaches to address performance, security, usability needs
 - Linux and Solaris kernels in kernel address space, so monolithic, plus modular for dynamic loading of functionality
 - Windows mostly monolithic, plus microkernel for different subsystem *personalities*
- Apple Mac OS X hybrid, layered, **Aqua** UI plus **Cocoa** programming environment
 - Below is kernel consisting of Mach microkernel and BSD Unix parts, plus I/O kit and dynamically loadable modules (called **kernel extensions**)



OPERATING-SYSTEM STRUCTURE

macOS and iOS Structure



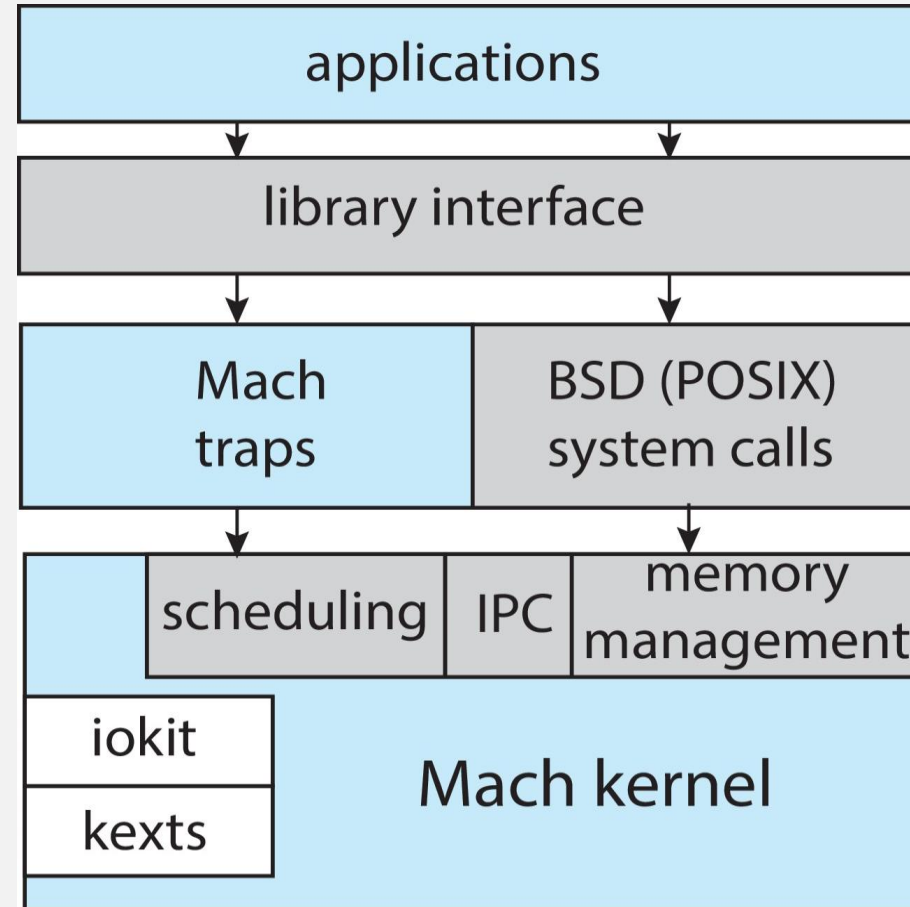
- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging





OPERATING-SYSTEM STRUCTURE

Darwin



- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging





OPERATING-SYSTEM STRUCTURE

iOS

- Apple mobile OS for *iPhone, iPad*
 - Structured on Mac OS X, added functionality
 - Does not run OS X applications natively
 - Also runs on different CPU architecture (ARM vs. Intel)
 - **Cocoa Touch** Objective-C API for developing apps
 - **Media services** layer for graphics, audio, video
 - **Core services** provides cloud computing, databases
 - **Core Operating System**, based on Mac OS X kernel

Cocoa Touch

Media Services

Core Services

Core OS

- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging





OPERATING-SYSTEM STRUCTURE

Android

- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging

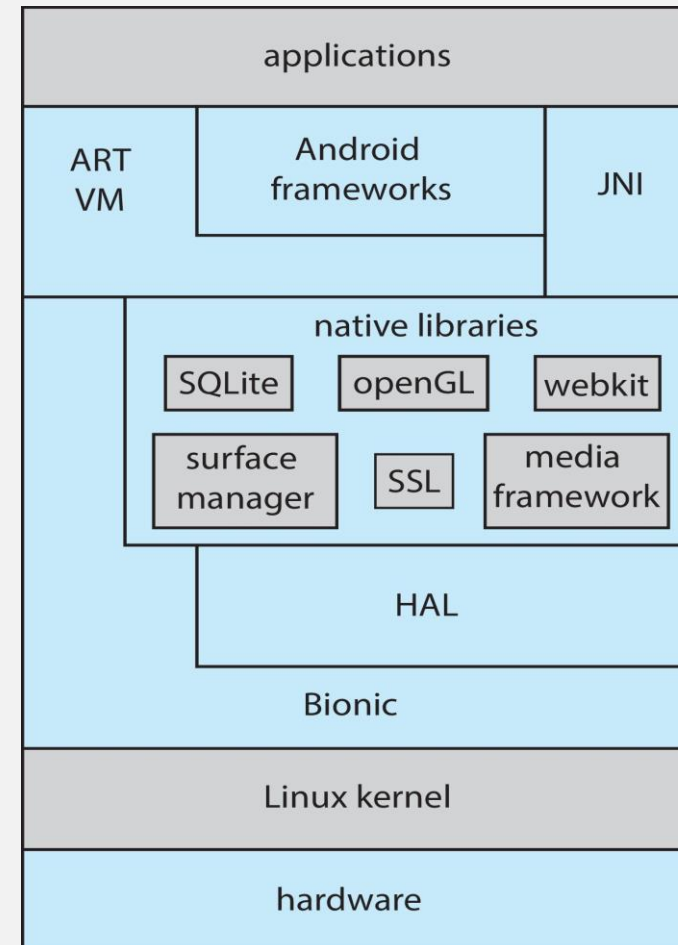


- Developed by Open Handset Alliance (mostly Google)
 - Open Source
- Similar stack to IOS
- Based on Linux kernel but modified
 - Provides process, memory, device-driver management
 - Adds power management
- **Runtime environment** includes core set of libraries and Dalvik virtual machine
 - Apps developed in Java plus Android API
 - ✓ Java class files compiled to Java bytecode then translated to executable than runs in Dalvik VM
- Libraries include frameworks for web browser (webkit), database (SQLite), multimedia, smaller libc



OPERATING-SYSTEM STRUCTURE

Android Architecture



- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- ✓ Operating-System Structure
- Building and Booting an Operating System
- Operating-System Debugging





BUILDING AND BOOTING AN OPERATING SYSTEM

- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- ✓ Building and Booting an Operating System
- Operating-System Debugging



- Operating Systems are generally designed to run on a class of systems with variety of peripherals
- Commonly, operating system is already installed on a purchased computer
 - But one can build and install some other operating systems
- If generating an operating system from scratch:
 - Write the operating system source code
 - Configure the OS for the system on which it will run
 - Compile the operating system
 - Install the operating system
 - Boot the computer and its new operating system



BUILDING AND BOOTING AN OPERATING SYSTEM

Building and Booting Linux

- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- ✓ Building and Booting an Operating System
- Operating-System Debugging



- Download Linux source code (<https://www.kernel.org>)
- Configure kernel via “**make menuconfig**”
- Compile the kernel using “**make**”
 - Produces **vmlinuz**, the kernel image
 - Compile kernel modules via “**make modules**”
 - Install kernel modules into **vmlinuz** via “**make modules_install**”
 - Install new kernel on the system via “**make install**”



BUILDING AND BOOTING AN OPERATING SYSTEM

System Boot

- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- ✓ Building and Booting an Operating System
- Operating-System Debugging



- When power is initialized on system, execution starts at a fixed memory location
- Operating system must be made available to hardware so hardware can start it
 - Small piece of code: **bootstrap loader, BIOS**, stored in **ROM** or **EEPROM** locates the kernel, loads it into memory, and starts it
 - Sometimes a two-step process where **boot block** at fixed location loaded by **ROM** code, **which loads bootstrap loader from disk**
 - Modern systems replace BIOS with **Unified Extensible Firmware Interface (UEFI)**
- Common bootstrap loader, **GRUB**, allows selection of kernel from multiple disks, versions, kernel options
- Kernel loads and system is then **running**
- Boot loaders frequently allow various boot states, such as single user mode



OPERATING-SYSTEM DEBUGGING

- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- ✓ Operating-System Debugging



- **Debugging** is finding and fixing **errors**, or **bugs**
- Can also include **performance tuning**
- OS generate **log files** containing error information
- Failure of an application can generate **core dump** file capturing memory of the process
- Operating system failure can generate **crash dump** file containing kernel memory



OPERATING-SYSTEM DEBUGGING

- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- ✓ Operating-System Debugging



- Beyond crashes, **performance tuning** can optimize system performance
 - Sometimes using *trace listings* of activities, recorded for analysis
 - **Profiling** is the periodic sampling of instruction pointer to look for statistical trends
- **Kernighan's Law:**
 - “Debugging is twice as hard as writing the code in the first place.
 - Therefore, if you write the code as cleverly as possible, you are, by definition, not smart enough to debug it.”



OPERATING-SYSTEM DEBUGGING

Performance Tuning

- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- ✓ Operating-System Debugging




- Improve performance by removing bottlenecks
- OS must provide means of computing and displaying measures of system behavior
- **Example:** “top” program or Windows Task Manager



OPERATING-SYSTEM DEBUGGING

Performance Tuning

The Windows 10 Task Manager



- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- ✓ Operating-System Debugging



OPERATING-SYSTEM DEBUGGING

Counters

- Outline
- Why Applications are
Operating-System Specific?
- Operating-System Design
and Implementation
- Operating-System Structure
- Building and Booting an
Operating System
- ✓ Operating-System
Debugging



- Operating systems **keep track of system activity** through a series of **counters**
 - the number of system calls made, or
 - the number of operations performed to a network device or disk
- **Counter-based tools** are tools that simply **inquire on the current value of certain statistics that are maintained by the kernel**



OPERATING-SYSTEM DEBUGGING

Counters

- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- ✓ Operating-System Debugging



■ Tools include:

- **ps** — reports information for a single process or selection of processes
- **top** — reports real-time statistics for current processes
- **vmstat** — reports memory-usage statistics
- **netstat** — reports statistics for network interfaces
- **iostat** — reports I/O usage for disks



- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- ✓ Operating-System Debugging



OPERATING-SYSTEM DEBUGGING

Tracing

- Collects data for a specific event, such as steps involved in a system call invocation
- Tools include:
 - **strace** – trace system calls invoked by a process
 - **gdb** – source-level debugger
 - **perf** – collection of Linux performance tools
 - **tcpdump** – collects network packets



OPERATING-SYSTEM DEBUGGING

BCC

- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- ✓ Operating-System Debugging



- **Debugging interactions** between **user-level** and **kernel code** nearly impossible without **toolset** that understands both and an instrument their actions
- **BCC (BPF Compiler Collection)** is a **rich toolkit** providing **tracing features for Linux**
- **Example: disksnoop.py** traces disk I/O activity

TIME(s)	T	BYTES	LAT(ms)
1946.29186700	R	8	0.27
1946.33965000	R	8	0.26
1948.34585000	W	8192	0.96
1950.43251000	R	4096	0.56
1951.74121000	R	4096	0.35

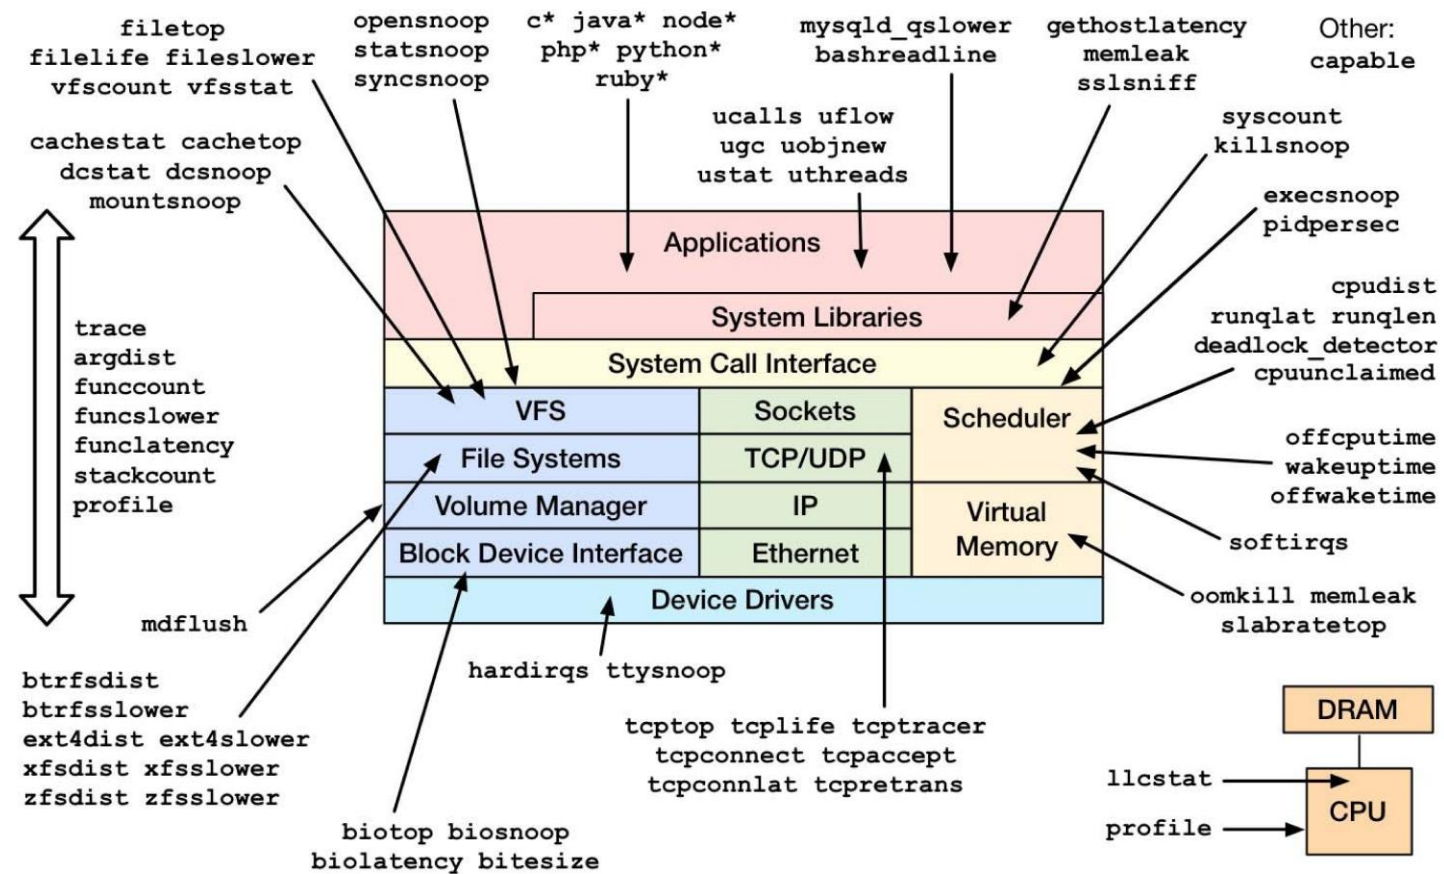


OPERATING-SYSTEM DEBUGGING

- Outline
- Why Applications are Operating-System Specific?
- Operating-System Design and Implementation
- Operating-System Structure
- Building and Booting an Operating System
- ✓ Operating-System Debugging



Linux bcc/BPF Tracing Tools



<https://github.com/iovisor/bcc#tools> 2017



2

CS 3104 OPERATING SYSTEMS

*** END ***

THANK YOU!

